

PROJECT REPORT

ON

SYSTEM MONITER TOOL

Name – Rimjhim Mohanty

Regno – 2241013268

Brach – Cse (Ds)

Institute Name - Institute of Technical Education and Research

Introduction

The System Monitor Tool is a Linux-based project developed in C++ to display real-time information about system processes, CPU usage, and memory utilization. Similar to the functionality of the top command, this tool collects data from the /proc file system and presents it in a clear and user-friendly terminal interface. It helps users monitor active processes, understand system resource consumption, and identify high-usage tasks. The project enhances understanding of Linux process management, CPU scheduling, and memory handling. It also provides hands-on experience with file handling, system calls, and real-time data updates in C++.

Objectives

- To develop a C++-based system monitoring application that displays real-time information such as running processes, CPU usage, and memory usage.
- To understand and utilize the Linux /proc filesystem for extracting process and system-level details like PID, process name, CPU time, and RAM usage.
- To implement features similar to the top command, including sorting processes based on CPU or memory usage and refreshing data at regular intervals.
- To improve knowledge of system programming, process management, and resource monitoring in Linux-based operating systems.
- To provide a user-friendly terminal interface for efficient system performance monitoring and analysis.

System Requirements

1. Hardware Requirements

- Minimum **2 GB RAM**
- **Dual-core processor** or higher
- At least **200 MB of free disk space**
- A system capable of running **Linux (Ubuntu or similar)**

2. Software Requirements

- **Operating System:** Ubuntu/Linux (recommended: Ubuntu 20.04 or later)

- **Compiler:** g++ with **C++17 support**
- **Libraries Used:**
 - <filesystem> (for process directory access)
 - Standard C++ Libraries (iostream, fstream, unistd.h, etc.)
- **Tools Required:**
 - Terminal / Command Line
 - Text Editor (VS Code, gedit, nano, etc.)
- **Access to /proc filesystem** for process and system data

Source Code / Implementation

```
#include <bits/stdc++.h>
#include <filesystem>
#include <unistd.h>
using namespace std;

struct Proc {
    int pid;
    string name;
    long rss_kb;
    unsigned long long cpu_jiffies;
    double cpu;
};

long long read_total_jiffies() {
    ifstream f("/proc/stat");
    string tag;
    long long a,b,c,d,e,fv,g,h,i,j;
```

```

f >> tag >> a >> b >> c >> d >> e >> fv >> g >> h >> i >> j;
return a + b + c + d + e + fv + g + h + i + j;
}

long long read_idle_jiffies() {
ifstream f("/proc/stat");
string tag;
long long user, nice, system, idle, iowait, irq, softirq, steal, guest, guest_nice;
f >> tag >> user >> nice >> system >> idle >> iowait >> irq >> softirq >> steal >>
guest >> guest_nice;
return idle + iowait;
}

bool read_stat(int pid, string &name, unsigned long long &ut_st, long &rss_kb) {
ifstream st("/proc/" + to_string(pid) + "/stat");
if (!st) return false;
string line;
getline(st, line);
auto lp = line.find('(');
auto rp = line.rfind(')');
if (lp == string::npos || rp == string::npos || rp <= lp) return false;
name = line.substr(lp + 1, rp - lp - 1);
string data = line.substr(rp + 2);
istringstream ss(data);
unsigned long long ut = 0, stt = 0;
long rss_pages = 0;
string skip;
for (int i = 0; i < 11; i++) ss >> skip;
ss >> ut >> stt;

```

```
for (int i = 0; i < 7; i++) ss >> skip;
ss >> rss_pages;
ut_st = ut + stt;
long page_kb = sysconf(_SC_PAGESIZE) / 1024;
rss_kb = rss_pages * page_kb;
return true;
}

vector<Proc> snapshot() {
vector<Proc> v;
for (auto &entry : filesystem::directory_iterator("/proc")) {
string name = entry.path().filename().string();
if (!all_of(name.begin(), name.end(), ::isdigit)) continue;
int pid = stoi(name);
string pname;
unsigned long long cpu;
long rss;
if (read_stat(pid, pname, cpu, rss)) {
v.push_back({pid, pname, rss, cpu, 0.0});
}
}
return v;
}

int main() {
int ncpu = max(1L, sysconf(_SC_NPROCESSORS_ONLN));
while (true) {
auto snap1 = snapshot();
long long t1 = read_total_jiffies();
```

```

long long idle1 = read_idle_jiffies();
usleep(1500000);
auto snap2 = snapshot();
long long t2 = read_total_jiffies();
long long idle2 = read_idle_jiffies();
unordered_map<int, Proc> map1, map2;
for (auto &p : snap1) map1[p.pid] = p;
for (auto &p : snap2) map2[p.pid] = p;
long long total_delta = t2 - t1;
double cpu_total = 100.0 * (double)(total_delta - (idle2 - idle1)) / total_delta;
vector<Proc> result;
for (auto &[pid, p2] : map2) {
    if (map1.find(pid) == map1.end()) continue;
    unsigned long long delta = p2.cpu_jiffies - map1[pid].cpu_jiffies;
    double cpu_usage = (double) delta * 100.0 * ncpu / total_delta;
    result.push_back({pid, p2.name, p2.rss_kb, p2.cpu_jiffies, cpu_usage});
}
sort(result.begin(), result.end(), [](Proc a, Proc b) {
    if (a.rss_kb != b.rss_kb) return a.rss_kb > b.rss_kb;
    return a.cpu > b.cpu;
});
cout << "\033[H\033[J";
cout << "CPU Usage: " << fixed << setprecision(2) << cpu_total << "%\n";
cout << left << setw(8) << "PID" << setw(25) << "Process"
<< setw(12) << "Memory(KB)" << setw(8) << "CPU%\n";
for (int i = 0; i < min((int)result.size(), 20); i++) {

```

```

cout << left << setw(8) << result[i].pid
<< setw(25) << result[i].name.substr(0, 24)
<< setw(12) << result[i].rss_kb
    << setw(8) << fixed << setprecision(2) << result[i].cpu << endl;
}

}

}

```

The screenshot shows a terminal window titled "proc_monitor.cpp" being edited in the nano 6.2 text editor. The code is a C++ program that reads CPU usage and process status from /proc/stat and /proc/[pid]/stat. It defines a struct Proc and two static inline functions to read total jiffies and idle jiffies. It then reads a process's stat line and extracts its name and RSS value. The code includes comments explaining the extraction logic and tokenization.

```

GNU nano 6.2
#include <bits/stdc++.h>
#include <filesystem>
#include <unistd.h>

using namespace std;

struct Proc {
    int pid;
    string name;
    long rss_kb;
    unsigned long long cpu_jiffies;
    double cpu; // computed percent
};

// ---- CPU totals from /proc/stat ----
static inline long long read_total_jiffies() {
    ifstream f("/proc/stat");
    string tag;
    long long user, nice, system, idle, iowait, irq, softirq, steal, guest, guest_nice;
    if (!(f >> tag >> user >> nice >> system >> idle >> iowait >> irq >> softirq >> steal >> guest >> guest_nice))
        return 0;
    return user + nice + system + idle + iowait + irq + softirq + steal + guest + guest_nice;
}

static inline long long read_idle_jiffies() {
    ifstream f("/proc/stat");
    string tag;
    long long user, nice, system, idle, iowait, irq, softirq, steal, guest, guest_nice;
    if (!(f >> tag >> user >> nice >> system >> idle >> iowait >> irq >> softirq >> steal >> guest >> guest_nice))
        return 0;
    return idle + iowait;
}

// ---- Per-process read from /proc/[pid]/stat and /status ----
static bool read_stat(int pid, string& name, unsigned long long& ut_st, long& rss_kb) {
    // Read /proc/[pid]/stat (one line)
    ifstream st("/proc/" + to_string(pid) + "/stat");
    if (!st) return false;

    string line;
    getline(st, line);
    if (line.empty()) return false;

    // Extract name between the last '(' and the last ')'
    size_t lp = line.find('(');
    size_t rp = line.rfind(')');
    if (lp == string::npos || rp == string::npos || rp <= lp) return false;
    name = line.substr(lp + 1, rp - lp - 1);

    // Tokenize after ")"
    string after = line.substr(rp + 2);
    vector<string> tok;
    tok.reserve(64);
    {
        istringstream ss(after);
        string w;
        while (ss >> w) tok.push_back(w);
    }
    // After ')', tokens start at field #3 (state). Need at least up to rss (field #24 -> index 21 here).
    if (tok.size() < 22) return false;
}

```

```
GNU nano 6.2                                         proc_monitor.cpp
    interval_sec = max(0.05, atof(argv[i + 1]));
}

const long ncpu = max(1L, sysconf(_SC_NPROCESSORS_ONLN));

while (true) {
    // First snapshot
    auto v1 = snapshot();
    long long tot1 = read_total_jiffies();
    long long idle1 = read_idle_jiffies();

    // Sleep for interval
    useconds_t us = static_cast<useconds_t>(interval_sec * 1e6);
    usleep(us);

    // Second snapshot
    auto v2 = snapshot();
    long long tot2 = read_total_jiffies();
    long long idle2 = read_idle_jiffies();

    // Build maps for delta
    unordered_map<int, Proc> m1, m2;
    m1.reserve(v1.size()); m2.reserve(v2.size());
    for (auto& p : v1) m1.emplace(p.pid, p);
    for (auto& p : v2) m2.emplace(p.pid, p);

    long long totald = max(1LL, tot2 - tot1);
    long long idled = max(0LL, idle2 - idle1);
    double cpu_total_active = 100.0 * (double)(totald - idled) / (double)totald; // 0..100 (all CPUs)

    vector<Proc> out;
    out.reserve(m2.size());
    for (const auto& kv : m2) {
        int pid = kv.first;
        const Proc& p2 = kv.second;

        auto it = m1.find(pid);
        if (it == m1.end()) continue; // skip brand-new processes (no baseline)

        unsigned long long dproc = 0;
        if (p2.cpu_jiffies >= it->second.cpu_jiffies)
            dproc = p2.cpu_jiffies - it->second.cpu_jiffies;

        // Scale to single-CPU 0..100% semantics
        double cpu = 100.0 * (double)ncpu * (double)dproc / (double)totald;
        out.push_back(Proc{p2.pid, p2.name, p2.rss_kb, p2.cpu_jiffies, cpu});
    }

    // Sort: RSS desc, then CPU desc
    sort(out.begin(), out.end(), [](const Proc& a, const Proc& b) {
        if (a.rss_kb != b.rss_kb) return a.rss_kb > b.rss_kb;
        return a.cpu > b.cpu;
    });

    clear_screen();
    cout << "Interval " << fixed << setprecision(2) << interval_sec
        << "s | CPUs: " << ncpu
        << " | Active CPU " << fixed << setprecision(2) << cpu_total_active << "%\n";
}
```

```

GNU nano 6.2
    unsigned long long ut = 0, stt = 0;
    long rss_pages = 0;
    try {
        ut = stoull(tok[11]);
        stt = stoull(tok[12]);
        rss_pages = stol(tok[21]);
    } catch (...) {
        return false;
    }
    ut_st = ut + stt;

    // Prefer VmRSS from /proc/[pid]/status (in kB)
    rss_kb = 0;
    ifstream sm("/proc/" + to_string(pid) + "/status");
    if (sm) {
        string k;
        while (sm >> k) {
            if (k == "VmRSS:") {
                long v;
                sm >> v;
                rss_kb = v; // already in kB
                break;
            }
        }
        if (rss_kb == 0) {
            long page_kb = sysconf(_SC_PAGESIZE) / 1024;
            rss_kb = rss_pages * page_kb;
        }
    }
    return true;
}

static vector<Proc> snapshot() {
    vector<Proc> v;
    for (auto& e : filesystem::directory_iterator("/proc")) {
        const string s = e.path().filename().string();
        if (s.empty() || !all_of(s.begin(), s.end(), ::isdigit)) continue;
        int pid = stoi(s);
        string name; unsigned long long utst; long rss;
        if (!read_stat(pid, name, utst, rss)) continue;
        v.push_back(Proc{pid, name, rss, utst, 0.0});
    }
    return v;
}

static void clear_screen() {
    // ANSI clear + home
    cout << "\033[H\033[J";
}

int main(int argc, char** argv) {
    ios::sync_with_stdio(false);
    cin.tie(nullptr);

    // Parse optional: -i <seconds>
    double interval_sec = 1.5;
    for (int i = 1; i + 1 < argc; ++i) {
        string a = argv[i];
        if ((a == "-i" || a == "--interval")) {

```

```

        cout << left << setw(8) << "PID"
        << setw(30) << "NAME"
        << setw(12) << "RSS(KB)"
        << setw(8) << "CPU%" << "\n";

        size_t rows = min<size_t>(30, out.size());
        for (size_t i = 0; i < rows; ++i) {
            const auto& p = out[i];
            string nm = p.name.size() > 28 ? p.name.substr(0, 28) : p.name;
            cout << left << setw(8) << p.pid
            << setw(30) << nm
            << setw(12) << p.rss_kb
            << setw(8) << fixed << setprecision(2) << p.cpu
            << "\n";
        }

        cout.flush();
    }
    return 0;
}

```

OUTPUT

Successful Compilation of the System Monitor Program

```
rimjhimmohanty@Rimjhim: ~  × + ▾ - □ ×
rimjhimmohanty@Rimjhim:~$ cd project3
rimjhimmohanty@Rimjhim:~/project3$ nano proc_monitor.cpp
rimjhimmohanty@Rimjhim:~/project3$ g++ proc_monitor.cpp -o proc_monitor -std=c++17
rimjhimmohanty@Rimjhim:~/project3$ ./proc_monitor
```

Execution of the System Monitor Tool

Interval 1.50s CPUs: 8 Active CPU 0.16%			
PID	NAME	RSS(KB)	CPU%
357	python3.10	86176	0.00
758	snapd	40184	0.00
309	unattended-upgr	20864	0.00
1609	packagekitd	20096	0.00
267	networkd-dispat	18304	0.00
62	systemd-journal	18184	0.00
206	systemd-resolve	14036	0.00
135	snapfuse	13348	0.00
1	systemd	12032	0.00
749	snapfuse	11316	0.00
100	snapfuse	10988	0.00
109	snapfuse	10512	0.00
161	snapfuse	9968	0.00
457	systemd	9472	0.00
149	snapfuse	9168	0.00
123	snapfuse	7752	0.00
1613	polkitd	7296	0.00
272	systemd-logind	7168	0.00
208	systemd-timesyn	7168	0.00
140	snapfuse	5920	0.00
87	systemd-udevd	5632	0.00

Viewing the /proc Directory in Linux

```
rimjhimmohanty@Rimjhim:~/project3$ ls /proc
1    208  471      driver      latency_stats  stat
100   260  62      dynamic_debug  loadavg      swaps
104   262  749     execdomains  locks        sys
109   267  758     fb          meminfo      sysrq-trigger
117   268  8       filesystems  misc        sysvipc
123   269  87      fs          modules     thread-self
1260  272  acpi     interrupts  mounts     timer_list
130   300  buddyinfo  iomem      mpt        tty
135   303  bus       ioports    mtrr      uptime
140   309  cgroups   irq        net        version
149   357  cmdline   kallsyms  pagetypeinfo vmallocinfo
154   4126 config.gz kcore      partitions  vmstat
160   4127 consoles  key-users  pressure   zoneinfo
1609  413  cpuinfo   keys      schedstat
161   4133 crypto    kmsq      scsi
1613  457  devices   kpagecgroup self
2    458  diskstats kpagecount slabinfo
206   4606 dma      kpageflags softirqs
rimjhimmohanty@Rimjhim:~/project3$ |
```

Conclusion

The System Monitor Tool effectively tracks real-time CPU and memory usage of running processes using data from the Linux /proc filesystem. It functions as a lightweight alternative to the top command while enhancing practical knowledge of system programming, process management, and C++ file handling. Overall, the project is efficient, educational, and provides a strong foundation for future enhancements.